

Thesis Title

A thesis submitted

in Partial Fulfillment of the Requirements
for the Degree of

Master of Technology

by

Anshu Avinash

to the

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING
INDIAN INSTITUTE OF TECHNOLOGY KANPUR

May, 2015

CERTIFICATE

It is certified that the work contained in the thesis titled **Thesis Title**, by **Anshu Avinash**, has been carried out under my supervision and that this work has not been submitted elsewhere for a degree.

Prof Piyush Kurur

Department of Computer Science & Engineering

IIT Kanpur

May, 2015

ABSTRACT

Name of student: **Anshu Avinash** Roll no: **10327122**

Degree for which submitted: **Master of Technology**

Department: **Computer Science & Engineering**

Thesis title: **Thesis Title**

Name of Thesis Supervisor: **Prof Piyush Kurur**

Month and year of thesis submission: **May, 2015**

Placeholder

Placeholder

Acknowledgements

Placeholder

Contents

List of Tables	xiii
List of Figures	xv
1 Introduction	1
1.1 Organization of the thesis	2
2 Related Work	3
2.1 Storing large objects in database	3
2.2 Storing metadata and filename in database	4
2.3 Comparison of both approaches	4
2.3.1 Performance	4
2.3.2 Security	4
3 Design	5
3.1 Creating a Blob	5
3.2 Writing to a Blob	5
3.3 Reading from Blob	6
3.4 Garbage Collection	6
4 Implementation	7
4.1 Functional Programming	7
5 Conclusion	9

References

List of Tables

3.1	Interface for operations on blob	6
3.2	Interface for garbage collection	6

List of Figures

Chapter 1

Introduction

Database refers to a collection of information that exists over a long period. A Database Management System (DBMS) is a tool for creating and managing large amount of data efficiently. Early database management systems evolved from file systems. These database systems used tree-based and the graph-based models for describing the structure of the information in a database.

Database systems changed significantly following a famous paper written by Ted Codd in 1970 [1]. Codd proposed that database systems should present the user with a view of data organized as tables called relations. This paper was foundation for popular relational databases like MySQL and PostgreSQL.

However, many of the web applications do not require the complex querying and management functionality offered by a Relational Database Management System (RDBMS). This among other reasons gave rise to NoSQL (Not only SQL) databases. These databases can be classified based on the data models used by them. Amazon's Dynamo [2] is a key-value store. In key-value stores records are stored and retrieved using a key that uniquely identifies the record. MongoDB [3] on the other hand is a document-oriented database. Document-oriented databases are designed for managing document-oriented information, also known as semi-structured data.

Today's web application also work with large files like images, music, videos etc. Size of these files can vary from few MBs to tens of GBs. The application developer can decide to store these files directly into the one of the databases mentioned above

or store it as a file and save the filename in the database.

In this thesis, we explore the second option and provide a simple interface written in *Haskell* to store large files. We also provide an interface for garbage collection of deleted blobs. We try to provide concurrency without using locks as much as possible.

1.1 Organization of the thesis

Chapter 2 discusses the approach of storing large files in databases. It also provides a background for this thesis work. In Chapter 3, we present our design. Chapter 4 describes our implementation. We conclude and present the future work in Chapter 5.

Chapter 2

Related Work

Large object files can either be stored directly in a database or we can store the path to the binary file and other metadata. In this section we will discuss few examples of both. We will also discuss merits and demerits of both the approaches.

2.1 Storing large objects in database

Exodus was one of the first databases to support storage of large object files [4]. It used B+ tree index on byte position within the object plus a collection of leaf (data) blocks. Exodus allowed searching for a range of bytes, inserting a sequence of bytes at a given point in the object, appending a sequence of bytes at the end of the object and to delete a sequence of bytes from a given point in the object.

Popular relational databases like MySQL and PostgreSQL both provide data types to store large object files. In MySQL the data type is called BLOB, and has operations similar to that on a string. Corresponding data type in PostgreSQL is `bytea`.

PostgreSQL also provides a BLOB data type which is quite different from MySQL's BLOB data type. Its implementation breaks large objects up into "chunks" and stores the chunks in rows in the database. A B-tree index guarantees fast searches for the correct chunk number when doing random access reads and writes.

A similar idea is used by MongoDB, which is a document database. It also divides

the large object into “chunks”. It uses GridFS specification for this [5]. GridFS works by storing the information about the file (called metadata) in the files collection. The data itself is broken down into pieces called chunks that are stored in the chunks collection.

2.2 Storing metadata and filename in database

Another approach to store large objects is to store only the filename and some metadata in the database. In this case the application has to take care of the all externally attached files as well as the security settings.

2.3 Comparison of both approaches

Both the approaches have their own benefits and disadvantages.

2.3.1 Performance

When we just store the filename in database, we skip the database layer altogether during file read and write operations. In the paper To BLOB or Not To BLOB [6], performance of SQL Server and NTFS has been compared. The results showed that the database gave higher throughputs for objects for relatively small size ($< 1\text{MB}$).

2.3.2 Security

Security and access controls are simplified when the data is directly stored in the database. When accessing the files directly, security settings between file system and database are independent from each other.

Chapter 3

Design

We store all the large objects in separate files. Our design is inspired from the maildir format [7]. All the large objects of a database are stored under a single directory which we also call a “BlobStore”. The BlobStore contains three subdirectories: *tmp*, *curr* and *old*. We will discuss purpose of this directories later in this chapter.

3.1 Creating a Blob

We provide a method called **createBlob** for creating a new blob. It takes a BlobStore as a parameter and returns a WriteContext. WriteContext contains the file handle of just created blob among other things. All the new blobs are created in the *tmp* folder. We use UUID-4 to give unique names to the newly created blobs.

3.2 Writing to a Blob

We only allow to add new data to the end of a given blob. We provide **writePartial** method for this. **writePartial** takes a blob and a WriteContext as arguments and appends the given blob to the WriteContext’s blob. Once all the data has been written to the blob, **finalizeWrite** is called. **finalizeWrite** takes a WriteContext as argument and moves the blob from *tmp* folder to *curr* folder. We also rename the file to SHA-512 hash of it’s contents. **finalizeWrite** returns a BlobId. This BlobId

contains the location of the blob.

3.3 Reading from Blob

Reading is also sequential. First the `initRead` method is called which returns a `ReadContext`, similar to the `WriteContext`. `ReadContext` also contains the file handle of the blob which is opened in read mode. `readPartial` takes a `ReadContext` and number of bytes as input and returns those number of bytes from the blob.

Table 3.1: Interface for operations on blob

Methods	Purpose
<code>createBlob</code>	Creates a blob in the given <code>BlobStore</code>
<code>writePartial</code>	Takes a blob and appends it to the end of the blob given in the argument
<code>finalizeWrite</code>	Takes a <code>WriteContext</code> as input and returns a <code>BlobId</code>
<code>initRead</code>	Takes a <code>BlobId</code> as input and returns a <code>ReadContext</code>
<code>readPartial</code>	Reads a given number of bytes from a <code>Blob</code>
<code>finalizeRead</code>	Completes the read

3.4 Garbage Collection

Table 3.2: Interface for garbage collection

Methods	Purpose
<code>startGC</code>	Starts garbage collection for the given <code>BlobStore</code>
<code>markBlobAsAccessible</code>	Marks the given blob as accessible
<code>endGC</code>	Ends the garbage collection by removing all the unaccessible blobs

Chapter 4

Implementation

In this section we will describe our implementation.

4.1 Functional Programming

Chapter 5

Conclusion

References

- [1] Edgar F Codd. “A relational model of data for large shared data banks”. In: *Communications of the ACM* 13.6 (1970), pp. 377–387.
- [2] Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall, and Werner Vogels. “Dynamo: amazon’s highly available key-value store”. In: *ACM SIGOPS Operating Systems Review*. Vol. 41. 6. ACM. 2007, pp. 205–220.
- [3] Kristina Chodorow. *MongoDB: the definitive guide.* ” O’Reilly Media, Inc.”, 2013.
- [4] Michael J Carey, David J DeWitt, Joel E Richardson, and Eugene J Shekita. *Object and file management in the EXODUS extensible database system*. University of Wisconsin-Madison. Computer Sciences Department, 1986.
- [5] David Hows, Peter Membrey, and Eelco Plugge. “GridFS”. In: *MongoDB Basics*. Springer, 2014, pp. 101–115.
- [6] Russell Sears, Catharine Van Ingen, and Jim Gray. “To blob or not to blob: Large object storage in a database or a filesystem?” In: *arXiv preprint cs/0701168* (2007).
- [7] Daniel J Bernstein. *Using maildir format*. 1995.
- [8] DS Justin Sheehy and D Smith. “Bitcask. a log-structured hash table for fast key/value data”. In: *White paper, April* (2010).
- [9] Mendel Rosenblum and John K Ousterhout. “The design and implementation of a log-structured file system”. In: *ACM Transactions on Computer Systems (TOCS)* 10.1 (1992), pp. 26–52.